
Porting the Jacobi Code to OpenMP

Introduction to High Performance Computing COMS30005

Adam Matheson
Student Number: P 1757290

1 Introduction

The trend in modern HPC is clear: parallelization is the way forward. In every aspect, performance is being found through concurrent execution; be that via vectorisation, multi-core CPUs or GPGPU computing. Despite new challenges arising from this new paradigm, the benefits on offer make this new way of programming more than worth the effort.

To illustrate the potential performance gains on offer, what follows is a walkthrough of a successful attempt to parallelize optimised serial code running the Jacobi algorithm for solving a set of linear equations, by porting it to OpenMP.

1.1 Methodology

After making a change, the program was re-profiled and the new run-time for matrices of orders 500, 1000, 2000 and 4000 recorded, averaged and rounded to an appropriate accuracy across three runs on BlueCrystal. Times are given in the format W/X/Y/Z for each matrix order respectively and are in seconds. For convenience, where significant performance gains were made, improvements are described in terms of “ tX times faster” (where t is the speedup); but smaller improvements are described in terms of percentages.

1.2 Serial Optimisations

Prior to the parallelization described in the following sections, a number of basic changes and serial optimisations were explored. Specifically, these were (in order, performance gain in reference to times after previous optimisation):

1. Change compiler from GNU (gcc 4.8.5) to Intel (icc 18.0.0): 3X faster
2. Ensure data order and data access pattern are the same (column vs. row traversal): 3.5X faster
3. Change data type from `double` to `float`: 5%-10% faster
4. Enable `-O3` compiler optimisation flag: 15% faster
5. Vectorisation check using `-qopt-report`
6. Attempted loop fusion (failed: introduced high error rate)

These optimisations resulted in runtimes at submission of 0.148/1.107/7.879/72.155.

1.2.1 Further Serial Optimisations

In addition to the above, one further optimisation became apparent after submission and was dutifully included in the serial code before embarking on this parallelization endeavour. The branching code `if (row != col)` in the `run()` method was removed entirely and instead, a looping technique that avoids iterating on that particular matrix element altogether was employed.

2 Parallelisation

2.1 Baseline

To judge relative performance improvements, we require a baseline runtime. The serial code with the above optimisations and changes from the first submission (continuing with `icc` and `-O3` optimisation), including the one further optimisation, ran in 0.137/0.994/6.891/66.731.

2.2 Performance Analysis

Relative performance improvements are good, but how does a programmer know if he is really running fast code? What is “fast”? We require a way to measure absolute performance gains, against a theoretical maximum performance. We can use the concept of “operational intensity” to produce a few approximations of a theoretical maximum performance.

2.2.1 Computational Complexity

It is useful to first understand the computational complexity (CC) of the Jacobi code in terms of “big O” notation. The key part of the Jacobi algorithm processes a matrix, nesting a `for` loop for processing columns inside another `for` loop for each row. If each of these takes $O(n)$ where n is the matrix order, then the CC is $O(n^2)$. This means it scales exponentially with the size of the data set.

2.2.2 Operational Intensity

We can use the CC to help work out the operational intensity (OI). This is a measure that can be described as “operations per byte of memory traffic”. Whereas CC only accounts for compute cost, OI can describe the relationship between compute and memory cost.

Looking at the Jacobi code, we can also identify the memory operations required during matrix processing. Working through the solver code block systematically, there are 8 bytes loaded/stored, and there are 2 operations executed, giving an OI of 2/8 or 0.25.

2.2.3 STREAM Benchmark

Another way of measuring absolute performance is with the STREAM benchmark. This measures the sustainable memory bandwidth for four long vector operations on specific hardware. Knowing the result of this benchmark for BlueCrystal can tell us how close we are to achieving peak memory bandwidth.

Running STREAM on BlueCrystal results in ...

2.2.4 Roofline Model

Now the OI and peak memory bandwidth can be used in a roofline model. This graph provides a visual means of identifying optimal performance in terms of a trade-off between compute and memory cost. We can also identify where the current code is under the “roof” of the graph, to determine whether the code is compute-bound or memory-bound, helping to focus optimisation and parallelization efforts. Here is the roofline model for BlueCrystal:

...

Taking the Jacobi OI of 0.25, it is clear this figure lies way to the left of the x -axis and is extremely memory bandwidth-bound.

2.3 Adding Pragmas

Now on to the parallelisation itself. The most obvious thing to do initially is to turn the `for` loops into parallel `for` loops. This can be done with the `pragma` ...

Several loops are available for parallelisation: ...

2.4 Profiling

To check the results of adding these OpenMP pragmas, profiling is necessary to get a better idea of exactly what the code is doing. Tau is a good profiler for multi-threaded code. The Tau output for the code after a quick and dirty parallelisation is:

...
gperftools?

2.5 Cache coherency/false sharing

Reduction

2.6 Reprofileing

After the above cleaning up the parallelisation, the code was re-profiled and the runtimes were now W/X/Y/Z.

Now is a good time to mention the overhead of parallelisation. Clearly the run times have not been decreasing in a direct relationship with the increasing number of cores (i.e. when using 4 cores instead of 1, the code is not 4 times faster). This is because parallelisation of code introduces some overhead with regards to thread management. Because the code is now more complex and things like memory conflicts and thread lifecycles must be managed, the performance gain is not directly related to the increase in processing power.

2.7 Libraries, BLAS, NAG C

According to the “seven dwarves” paper (REFERENCE), the vast majority of code executed in HPC falls into one of seven categories. This concept is useful because knowing which of the “dwarves” a piece of code is, may mean that historically, similar code has been dealt with before. Looking at these past examples may provide help on how to better solve your problem.

The Jacobi solver falls into the “dense linear algebra” category. From research, the BLAS and NAG C libraries are applicable to this category of HPC. These libraries contain many common mathematical functions, written in a highly optimised manner - for example, matrix manipulation. These libraries were tested in the parallel code where appropriate, and the resulting times after swapping home-grown code for the library code was W/X/Y/Z.

2.8 Re-testing compiler

Compiler optimisations can be somewhat of a mystery. The complexity of modern compilers means their behaviour is not always predictable in advance. Sometimes, less aggressive optimisation flags can be faster than more aggressive flags, for example.

Testing of the following compiler flags was undertaken after parallelisation to check behaviour:

- Optimisation flags -O, -O2, -O3:
-

3 Conclusion

Overall,

162 **3.1 Going Further - Scaling**
163
164 One final useful thing to know is how well a piece of code will scale with further paralleliza-
165 tion. Testing the Jacobi algorithm on a growing number of cores results in the following
166 graph:
167
168 ...
169 Which means the scaling is type...
170 So can be scaled more/cannot be scaled that well
171
172 Amdahl's Law Gustavson's Law
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215