

---

# Porting the Jacobi Code to OpenMP

## Introduction to High Performance Computing COMS30005

---

Adam Matheson  
Student Number: P 1757290

### 1 Introduction

The trend in modern HPC is clear: parallelization is the way forward. In every aspect, performance is being found through concurrent execution; be that via vectorisation, multi-core CPUs or even GPGPU computing. Despite programming challenges arising from this new paradigm, the benefits on offer make it more than worth the effort.

To illustrate the potential performance gains on offer, what follows is a walkthrough of a successful attempt to parallelize optimised serial code running the Jacobi algorithm for solving a set of linear equations, by porting it to OpenMP.

#### 1.1 Methodology

After making a change, the program was re-profiled and the new run-time for matrices of orders 500, 1000, 2000 and 4000 recorded, averaged and rounded to an appropriate accuracy across three runs on BlueCrystal. Times are given in the format W/X/Y/Z for each matrix order respectively and are in seconds. For convenience, where significant performance gains were made, improvements are described in terms of “ $tX$  times faster” (where  $t$  is the speedup); whereas smaller improvements are described in terms of percentages.

#### 1.2 Serial Optimisations

Prior to parallelization, a number of basic changes and optimisations of the serial code were explored. In order of implementation, with performance gain relative to the preceding optimisation’s times, these were:

1. Change compiler from GNU (gcc 4.8.5) to Intel (icc 18.0.0): **3X faster**
2. Ensure data order and data access pattern are the same (column vs. row traversal): **3.5X faster**
3. Change data type from `double` to `float`: **5%-10% faster**
4. Enable `-O3` compiler optimisation flag: **15% faster**
5. Vectorisation check using `-qopt-report`

These optimisations resulted in runtimes at submission of 0.148/1.107/7.879/72.155 (down from approximately 1.5/10/130/1180 for unoptimised code).

##### 1.2.1 Further Serial Optimisations

In addition to the above, one further optimisation became apparent after submission and was dutifully included in the serial code before embarking on this parallelization endeavour. The branching code `if (row != col)` in the `run()` method was removed entirely and instead, a looping technique that avoids iterating on that particular matrix element altogether was employed, thus better facilitating pipelining and vectorisation.

## 2 Parallelisation

### 2.1 Baseline

To judge relative performance improvements, a baseline runtime is required. The serial code with the above optimisations and changes from the first submission (continuing with `icc` and `-O3` optimisation), including the one further optimisation, produced runtimes of 0.137/0.994/6.891/66.731.

### 2.2 Performance Analysis

Relative performance improvements are good, but how does a programmer know if he is really running fast code? What exactly is “fast”? Measuring absolute performance gains against a theoretical maximum performance can help answer this. The concept of “operational intensity” can be used to produce a few illustrations of theoretical maximum performance.

#### 2.2.1 Computational Complexity

It is useful to first understand the computational complexity (CC) of the Jacobi code in terms of “big O” notation. The key part of the Jacobi algorithm processes a matrix, nesting two adjacent `for` loops that process matrix columns inside another `for` loop, processing each row:

```
1 // Perform Jacobi iteration
2 for (row = 0; row < N; row++)
3 {
4     dot = 0.0;
5     skip_count = N;
6
7     for (col = 0; col < row; col++)
8     {
9         dot += A[row*N + col] * x[col];
10    }
11    // Skip matrix element where col==row
12    for (col = (row + 1); col < N; col++)
13    {
14        dot += A[row*N + col] * x[col];
15    }
16
17    xtmp[row] = (b[row] - dot) / A[row*N + row];
18 }
```

If each of these loops takes  $O(n)$  where  $n$  is the matrix order, then the CC is  $O(n^2)$ . This means the computational requirement scales exponentially with the size of the matrix order.

#### 2.2.2 Operational Intensity

Whereas CC only accounts for compute cost, operational intensity (OI) can describe the relationship between compute and memory cost. This is a measure that can be described as “operations per byte of memory traffic”.

By simply looking at the solver portion of the Jacobi code, we can identify the memory operations required during matrix processing. Working through the code block systematically, there are 8 bytes loaded/stored, and there are 2 operations executed, giving an OI of  $2/8$  or 0.25.

#### 2.2.3 STREAM Benchmark

Another way of measuring absolute performance is with the STREAM benchmark. This measures the peak sustainable memory bandwidth on specific hardware. Knowing the result

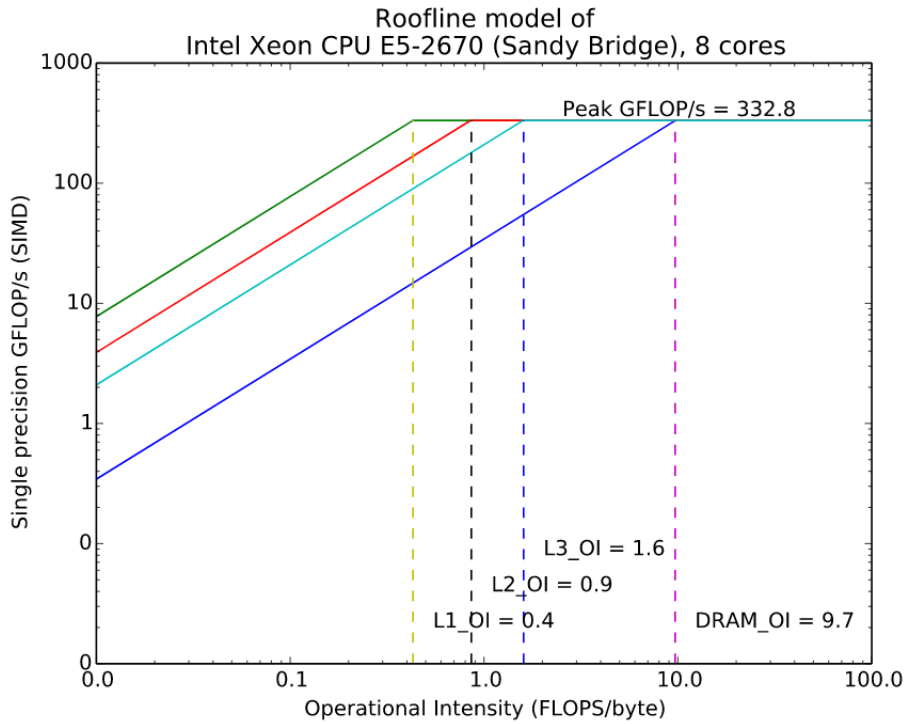


Figure 1: Cache-aware roofline model for a BCP3 CPU

of this benchmark for BlueCrystal can tell us how close we are to achieving peak memory bandwidth.

STREAM benchmark results available from Simon McIntosh-Smith for an Intel Xeon CPU E5-2670 (SB) 8-cores gave the following results:

Memory Level	Bandwidth in GB/s
L1 cache	776
L2 cache	390
L3 cache	209
DRAM	34.4

#### 2.2.4 Roofline Model

Now the OI and peak memory bandwidth can be used in a roofline model. This graph provides a visual means of identifying optimal performance in terms of a trade-off between compute and memory cost. We can also identify where the code is under the “roof” of the graph, to determine whether the code is compute-bound or memory-bound, helping to focus optimisation and parallelization efforts.

Figure 1 shows a cache-aware roofline model for a BlueCrystal CPU. Taking the Jacobi OI of 0.25, it is clear this figure lies way below the peak performance point with reference to DRAM. However it is not particularly far from the peak performance mark of the L1 cache. If code modifications took place and the Jacobi OI became above 0.4, and if the code was hitting the L1 cache the majority of the time, this figure would become relevant because then, the code would change to be compute bound!

#### 2.3 Adding Pragmas

Now on to the parallelisation itself. The most obvious thing to do initially is to turn the `for` loops into parallel `for` loops. This can be done with the pragma ...

Several loops are available for parallelisation: ...

## 2.4 Profiling

To check the results of adding these OpenMP pragmas, profiling is necessary to get a better idea of exactly what the code is doing. **Tau** is a good profiler for multi-threaded code. The **Tau** output for the code after a quick and dirty parallelisation is:

...  
gperftools?

## 2.5 Cache coherency/false sharing

Reduction

## 2.6 Reprofileing

After the above cleaning up the parallelisation, the code was re-profiled and the runtimes were now W/X/Y/Z.

Now is a good time to mention the overhead of parallelisation. Clearly the run times have not been decreasing in a direct relationship with the increasing number of cores (i.e. when using 4 cores instead of 1, the code is not 4 times faster). This is because parallelisation of code introduces some overhead with regards to thread management. Because the code is now more complex and things like memory conflicts and thread lifecycles must be managed, the performance gain is not directly related to the increase in processing power.

## 2.7 Libraries

According to the “seven dwarves” paper, the vast majority of code executed in HPC falls into one of seven categories. This concept is useful because knowing which of the “dwarves” a piece of code is, may mean that historically, similar code has been dealt with before. Looking at these past examples may provide help on how to better solve your problem.

The Jacobi solver falls into the “dense linear algebra” category. From research, the BLAS and NAG C libraries are applicable to this category of HPC. These libraries contain many common mathematical functions, written in a highly optimised manner - for example, matrix manipulation. These libraries were tested in the parallel code where appropriate, and the resulting times after swapping home-grown code for the library code was W/X/Y/Z.

## 2.8 Re-testing compiler

Compiler optimisations can be somewhat of a mystery. The complexity of modern compilers means their behaviour is not always predictable in advance. Sometimes, less aggressive optimisation flags can be faster than more aggressive flags, for example.

Testing of the following compiler flags was undertaken after parallelisation to check behaviour:

- Optimisation flags -O, -O2, -O3:
- 

## 3 Conclusion

Overall,

216  
217  
218  
219  
220  
221  
222  
223  
224  
225  
226  
227  
228  
229  
230  
231  
232  
233  
234  
235  
236  
237  
238  
239  
240  
241  
242  
243  
244  
245  
246  
247  
248  
249  
250  
251  
252  
253  
254  
255  
256  
257  
258  
259  
260  
261  
262  
263  
264  
265  
266  
267  
268  
269

### 3.1 Going Further

In terms of further optimisation, loop fusion was attempted in the Jacobi solver, to combine matrix processing with convergence check. However this introduced a high error rate (probably bug). Successfully implementing this could realise further performance gains.

One final useful thing to know is how well a piece of code will scale with further parallelization. Testing the Jacobi algorithm on a growing number of cores results in the following graph:

...

Which means the scaling is type...

So can be scaled more/cannot be scaled that well

Amdahl's Law Gustavson's Law