
Porting the Jacobi Code to OpenMP

Introduction to High Performance Computing COMS30005

Adam Matheson
Student Number: P 1757290

1 Introduction

The trend in modern HPC is clear: parallelization is the way forward. In every aspect, performance is being found through concurrent execution; be that via vectorisation, multi-core CPUs or even GPGPU computing. Despite programming challenges arising from this new paradigm, the benefits on offer make it more than worth the effort.

To illustrate the potential performance gains on offer, what follows is a walkthrough of a successful attempt to parallelize optimised serial code running the Jacobi algorithm for solving a set of linear equations, by porting it to OpenMP.

1.1 Methodology

After making a change, the program was re-profiled and the new run-time for matrices of orders 500, 1000, 2000 and 4000 recorded, averaged and rounded to an appropriate accuracy across three runs on BlueCrystal. Times are given in the format W/X/Y/Z for each matrix order respectively and are in seconds. For convenience, where significant performance gains were made, improvements are described in terms of “ tX times faster” (where t is the speedup); whereas smaller improvements are described in terms of percentages.

1.2 Serial Optimisations

Prior to parallelization, a number of basic changes and optimisations of the serial code were explored. In order of implementation, with performance gain relative to the preceding optimisation’s times, these were:

1. Change compiler from GNU (gcc 4.8.5) to Intel (icc 18.0.0): **3X faster**
2. Ensure data order and data access pattern are the same (column vs. row traversal):
3.5X faster
3. Change data type from `double` to `float`: **5%-10% faster**
4. Enable `-O3` compiler optimisation flag: **15% faster**
5. Vectorisation check using `-qopt-report`

These optimisations resulted in runtimes at submission of 0.148/1.107/7.879/72.155 (down from approximately 1.5/10/130/1180 for unoptimised code).

1.2.1 Further Serial Optimisations

In addition to the above, one further optimisation became apparent after submission and was dutifully included in the serial code before embarking on this parallelization endeavour. The branching code `if (row != col)` in the `run()` method was removed entirely and instead, a looping technique that avoids iterating on that particular matrix element altogether was employed, thus better facilitating pipelining and vectorisation.

2 Parallelisation

2.1 Baseline

To judge relative performance improvements, a baseline runtime is required. The serial code with the above optimisations and changes from the first submission (continuing with `icc` and `-O3` optimisation), including the one further optimisation, produced runtimes of 0.137/0.994/6.891/66.731.

2.2 Performance Analysis

Relative performance improvements are good, but how does a programmer know if he is really running fast code? What exactly is “fast”? Measuring absolute performance gains against a theoretical maximum performance can help answer this. The concept of “operational intensity” can be used to produce a few illustrations of theoretical maximum performance.

2.2.1 Computational Complexity

It is useful to first understand the computational complexity (CC) of the Jacobi code in terms of “big O” notation. The key part of the Jacobi algorithm (henceforth referred to as the “solver” code block”) processes a matrix, nesting two adjacent `for` loops that process matrix columns inside another `for` loop, processing each row:

```
1 // Perform Jacobi iteration
2 for (row = 0; row < N; row++)
3 {
4     dot = 0.0;
5     skip_count = N;
6
7     for (col = 0; col < row; col++)
8     {
9         dot += A[row*N + col] * x[col];
10    }
11    // Skip matrix element where col==row
12    for (col = (row + 1); col < N; col++)
13    {
14        dot += A[row*N + col] * x[col];
15    }
16
17    xtmp[row] = (b[row] - dot) / A[row*N + row];
18 }
```

If each of these loops takes $O(n)$ where n is the matrix order, then the CC is $O(n^2)$. This means the computational requirement scales exponentially with the size of the matrix order.

2.2.2 Operational Intensity

Whereas CC only accounts for compute cost, operational intensity (OI) can describe the relationship between compute and memory cost. This is a measure that can be described as “operations per byte of memory traffic”.

By simply looking at the solver portion of the Jacobi code, we can identify the memory operations required for the solver. Working through the solver block systematically, there are 8 bytes loaded/stored, and there are 2 operations executed, giving an OI of $2/8$ or 0.25.

2.2.3 STREAM Benchmark

Another way of measuring absolute performance is with the STREAM benchmark. This measures the peak sustainable memory bandwidth on specific hardware. Knowing the result of this benchmark for BlueCrystal can tell us how close we are to achieving peak memory bandwidth.




Figure 1: Cache-aware roofline model for a BCP3 CPU

STREAM benchmark results available from Simon McIntosh-Smith for an Intel Xeon CPU E5-2670 (SB) 8-cores gave the following results:

Memory Level	Bandwidth in GB/s
L1 cache	776
L2 cache	390
L3 cache	209
DRAM	34.4

2.2.4 Roofline Model

Now the OI and peak memory bandwidth can be used in a roofline model. This graph provides a visual means of identifying optimal performance in terms of a trade-off between compute and memory cost. We can also identify where the code is under the “roof” of the graph, to determine whether the code is compute-bound or memory-bound, helping to focus optimisation and parallelization efforts.

Figure 1 shows a cache-aware roofline model for a BlueCrystal CPU. Taking the Jacobi OI of 0.25, it is clear this figure lies way below the peak performance point with reference to DRAM. However it is not particularly far from the peak performance mark of the L1 cache. If code modifications took place and the Jacobi OI became above 0.4, and if the code was hitting the L1 cache the majority of the time, this figure would become relevant because then, the code would change to be compute bound!

2.3 Adding Pragmas

Now on to the parallelisation itself. The most obvious thing to do initially is to turn the hungriest `for` loops into parallel `for` loops. From previous profiling during serial optimisation, it is known that the outer loop in the solver code block occupies most of the runtime.

162 To parallelise this loop, `pragma omp parallel for shared(xtmp) private(dot,`
163 `skip_count)` is used.

164 Several loops are available for parallelisation, however

166 2.4 Profiling

168 To check the results of adding these OpenMP pragmas, profiling is necessary to get a better
169 idea of exactly what the code is doing. **Tau** is a good profiler for multi-threaded code. The
170 **Tau** output for the code after a quick and dirty parallelisation is:

171 ...
172 gperftools?

175 2.5 Cache coherency/false sharing

177 Reduction

179 2.6 Reprofile

181 After the above cleaning up the parallelisation, the code was re-profiled and the runtimes
182 were now W/X/Y/Z.

183 Now is a good time to mention the overhead of parallelisation. Clearly the run times have
184 not been decreasing in a direct relationship with the increasing number of cores (i.e. when
185 using 4 cores instead of 1, the code is not 4 times faster). This is because parallelisation
186 of code introduces some overhead with regards to thread management. Because the code is
187 now more complex and things like memory conflicts and thread lifecycles must be managed,
188 the performance gain is not directly related to the increase in processing power.

190 2.7 Libraries

192 According to the “seven dwarves” paper, the vast majority of code executed in HPC falls
193 into one of seven categories. This concept is useful because knowing which of the “dwarves”
194 a piece of code is, may mean that historically, similar code has been dealt with before.
195 Looking at these past examples may provide help on how to better solve your problem.

196 The Jacobi solver falls into the “dense linear algebra” category. From research, the BLAS
197 and NAG C libraries are applicable to this category of HPC. These libraries contain many
198 common mathematical functions, written in a highly optimised manner - for example, matrix
199 manipulation. These libraries were tested in the parallel code where appropriate, and the
200 resulting times after swapping home-grown code for the library code was W/X/Y/Z.

202 2.8 Re-testing compiler

204 Compiler optimisations can be somewhat of a mystery. The complexity of modern compilers
205 means their behaviour is not always predictable in advance. Sometimes, less aggressive
206 optimisation flags can be faster than more aggressive flags, for example.

207 Testing of the following compiler flags was undertaken after parallelisation to check be-
208 haviour:

- 209 • Optimisation flags -O, -O2, -O3:
- 211 •

213 3 Conclusion

215 Overall,

216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269

3.1 Going Further

In terms of further optimisation, loop fusion was attempted in the Jacobi solver, to combine matrix processing with convergence check. However this introduced a high error rate (probably bug). Successfully implementing this could realise further performance gains.

One final useful thing to know is how well a piece of code will scale with further parallelization. Testing the Jacobi algorithm on a growing number of cores results in the following graph:

...

Which means the scaling is type...

So can be scaled more/cannot be scaled that well

Amdahl's Law Gustavson's Law